

SHIVAM KUMAR VERMA

shivam.verma.dev00@gmail.com • +91-9135668614 • [Website](#) • [LinkedIn](#) • [GitHub](#)

PROFESSIONAL SUMMARY

Highly motivated Software Developer with a strong foundation in full-stack development, specializing in building scalable, production-grade applications using modern technologies like TypeScript, React, Node.js, and Python. Possesses hands-on experience with LLM integration, prompt engineering, and optimizing AI-driven features through an impactful internship and complex personal projects. Proven ability to design, develop, and debug robust software solutions. Eager to apply strong problem-solving skills and competitive programming background to evaluate and refine AI-generated code and contribute to advancing large language models at Turing.

SKILLS

Languages: Python, JavaScript, TypeScript, C++

Frameworks: React, Node.js, Express.js, Next.js, FastAPI

AI & ML: LLM integration, Prompt engineering, Semantic search, Pinecone

Databases & ORMs: PostgreSQL, Redis, ClickHouse, Prisma, Drizzle

DevOps & Cloud: Docker, AWS (EC2, S3, Lambda), Nginx, Linux, Git, CI/CD

EXPERIENCE

Software Developer Intern at klimb.io

Nov 2025 – Current •

- Engineered a credit-based monetization system, migrating platform access from feature flags to a robust PostgreSQL-backed model with middleware guards, enabling monetization of 6+ AI features and reducing system complexity.
- Implemented a global field template architecture by restructuring the database schema and backfilling data, achieving zero-downtime migration with complete data integrity across all customer tenants.
- Optimized recruiter dashboard performance by introducing precomputed, indexed data pipelines, significantly reducing load times from 2-3 seconds to under 500ms.

PROJECTS

CodeSM [React](#), [Node.js](#), [Express.js](#), [Redis](#), [BullMQ](#), [Docker](#), [PostgreSQL](#), [AWS EC2](#)

- Developed a scalable asynchronous code execution engine utilizing segregated BullMQ queues and Redis payload caching to ensure reliable, high-throughput code submissions with near-zero dropped jobs.
- Implemented a persistent Docker container strategy with tmpfs mounts, reducing per-test execution time from 300ms to 5ms (98% faster) by eliminating container spin-up overhead.
- Designed a fault-tolerant worker architecture incorporating a Dead Letter Queue for retries and a 1-hour Redis cache layer, significantly reducing database load during high-traffic periods.

Storix [Next.js](#), [TypeScript](#), [FastAPI](#), [Celery](#), [Redis](#), [PostgreSQL](#), [AWS S3](#)

- Engineered an asynchronous PDF processing pipeline using distributed Celery workers and a Redis task queue, enabling non-blocking long-running jobs without impacting API response times.
- Implemented direct cloud uploads by generating pre-signed AWS S3 URLs for client-side transfers, effectively offloading file transfer from the server and improving upload reliability.
- Integrated real-time job status updates via a Pub/Sub and WebSocket layer, providing users with instant feedback from file upload through delivery.

Chatterly [Next.js](#), [Socket.IO](#), [Redis](#), [PostgreSQL](#), [Prisma](#), [NextAuth](#)

- Built a horizontally scalable messaging system by leveraging Redis Pub/Sub for cross-server event propagation, ensuring consistent room state across multiple distributed instances.
- Designed zero-drop room event routing by subscribing each server to dedicated Redis channels, guaranteeing message delivery under concurrent multi-user loads.
- Secured private rooms by implementing JWT verification and room-level authorization checks within the Socket.IO middleware, preventing unauthorized access.

EDUCATION

B.Tech in Computer Science and Engineering from Indian Institute of Information Technology Nagpur Nov 2022 – May 2026 • Nagpur, Maharashtra

CGPA: 7.5 | Relevant Coursework: Operating Systems, Database Management Systems, Computer Networks, Object-Oriented Programming